

Show Dossier — Redesign & Data Specification

Information architecture for the LosslessBob show dossier, the data each anchor consumes, and the fields that must be generated before the redesign can ship.

Rev A · draft 9 September 2026 Reference mock: Show Dossier — Redesign (2010-03-29).html

Baseline: dossier template @ master 2026-08-20_003101

- 1 Scope & goals
- 2 Section order & rationale
- 3 Provenance tiers
- 4 Audit of the reference mock
- 5 New data requirements
- 6 Anchor reference (exhaustive)
- 7 Defects in the current generator
- 8 Rendering rules & degradation
- 9 Open questions

1 Scope & goals

The dossier answers one question — **which copy of this show should I get, and why** — and provides the reference apparatus to justify the answer. The current template presents the reference apparatus first and the answer fourth, gives every source equal visual weight regardless of how much it differs from its neighbours, and passes several scraped text blobs through to the page unparsed.

This spec covers three things: the reordered section sequence and the reasoning behind it; every anchor point in the new layout with the field that populates it; and the data that does not exist yet and must be generated, acquired, or curated before the corresponding anchor can render. It does not cover visual styling, which is inherited from the existing token set with one addition (a reserved warm accent used exclusively for the recommended copy).

Two hard constraints follow from the audit in §4. First, **no anchor may render a value the generator cannot trace to a source** — the reference mock violates this in six places and those violations are the substance of §5. Second, **every field must have a defined null rendering**, because coverage across the catalog is wildly uneven: the same template serves a 1965 BBC broadcast with eighteen near-identical B+ transfers and no scan grades at all, and a 2010 audience-taped show with nine sources and full grading.

2 Section order & rationale

#	SECTION	CONTAINS	WHY HERE
1	Header	Date as H1, venue/city subtitle, tour + NET number + indices, sibling-night strip	A Dylan show's identity is its date, not its venue. The run strip answers "was this the good night?" before scrolling.
2	Verdict	Recommended copy, specs, why-this, comparison to runner-up, alternates by axis, score ledger (collapsed)	The page's reason for existing. Absorbs the old at-a-glance strip, which duplicated fields that appear in the sources table anyway.
3	Setlist	Songs with instrumentation, encore banding, premiere markers, stats rail, band roster	This is the show. It is also the only section a casual reader wants, and it was previously below an eighteen-row reference table.
4	Sources	Top 4 ranked, remainder collapsed, family banding, lineage footnote	Reference apparatus. Belongs after the answer and the content, before the surrounding context.
5	Context	Chronicle prose, bobtalk, venue card with run/city counts	Editorial framing; read last or not at all. Bobtalk is promoted into §3 when it exists and can be attached to songs.
6	Related	City/venue history, cross-reference links	Exit routes. Replaces the location map, which consumed 300 px of prime space to convey the word "Tokyo".
7	Provenance	Credits, generation stamps, local-analysis disclosure	Unchanged from baseline; already correct.

DEVIATION FROM THE MOCK

The mock renders the location card inside §5 (Context) rather than as its own section. Keep it there: the venue is context, and a standalone Location section cannot justify its own heading once the map is gone.

3 Provenance tiers

Every field in §6 carries a tier. The tier determines who is responsible for producing it, and — critically — whether the renderer is allowed to state it as fact.

T1 Available now

The current generator already emits this. No work beyond wiring it to a new anchor.

T2 Derivable

Computable from data already present in the dossier payload — arithmetic, parsing a string the generator already has, or grouping. No new source needed.

T3 Must be generated

Requires data the payload does not contain: a new query against the LB catalog, a new external source, or a new computation over the whole show corpus.

T4 Editorial

Requires human judgement or a composed sentence. May be templated from lower-tier fields but must be marked as local analysis.

Fields are emitted as objects, not bare values, so the renderer can honour the tier:

```
{ value: "24/96", tier: 2, source: "lineage_string:LB-08485", confidence: "certain" }
{ value: null, tier: 3, source: null, confidence: "unavailable" }
```

A field with `tier ≥ 3` and `confidence: "inferred"` renders with the inference marker defined in §8. A field with `value: null` renders its defined fallback and never renders an invented substitute.

4 Audit of the reference mock

The mock was built by hand from the 2010-03-29 dossier. Most of it is faithful; the exceptions are listed here in full because they define the work in §5. Anything not listed below came from the dossier payload verbatim.

4.1 Derived by hand, checkable, safe to automate

VALUE IN MOCK	DERIVATION	TIER
126 min · 80 + 46	Sum of the two half-set lengths already emitted per source.	T2
night 7 of 7	Count of Zepp Tokyo March-2010 dates in the venue-history blob (6) plus the current show.	T2
45 Tokyo shows	Sum of the venue-history blob, grouped by year and venue.	T2
24/96 · 16/44.1	Parsed out of the lineage strings for LB-08485 and LB-08493. Absent from the other seven, which correctly render "—".	T2
harp on 8 of 18	Cardinality of the harmonica song-range list in the lineup blob (2, 4–7, 10, 14, 18).	T2
guitar only on Baby Blue	From the "2 Bob Dylan (guitar)" clause; song 2 is <i>It's All Over Now, Baby Blue</i> .	T2
Dylan–Hunter etc.	Parsed from parenthetical composer credits inside the song titles.	T2
A– vs B+ is 5 points	84 – 79, both already emitted.	T2

4.2 Invented — no backing in the payload

MATERIAL FINDING

The three fields that make the new verdict card compelling — **completeness**, **tour premieres**, and **official release status** — are precisely the three with no data behind them. The layout must not ship until each has a real source or an agreed null rendering.

VALUE IN MOCK	PROBLEM	RESOLUTION
Completeness 18/18 + bar	No per-song presence data exists anywhere in the payload. Asserted from runtime plausibility (126 min for an 18-song set).	Req D–01
"new to tour" on songs 13 & 16	The payload states "2 new songs for this tour" but never which two. <i>Thunder On The Mountain</i> and <i>Jolene</i> were a guess — made while papering over the mis-mapped subnote described in §7.2.	Req D–02
Official release – none indexed	Both the field and its null value were invented. No release-status data is ingested today.	Req D–03
"biggest rotation of the Japan run"	Requires the rotation percentage for the other six nights; only this show's 72% is available.	Req D–04
Generation: Master	Defensible for the silver-disc case (the lineage says so) but "master" on LB-08485/08476/08537 is inferred from the shape of a mic→preamp→recorder chain, not stated.	Req D–05
Odaiba	Neighbourhood supplied from the author's own knowledge, not the payload.	Drop or Req D–06
"closing the 2010 Tour of Japan"	29 March is the last <i>Tokyo</i> date in the blob; the payload does not establish that the tour ended there.	Req D–07
"taper unknown"	Absence of a taper in the dossier is not the same as unknown in the catalog. Overstates a null as a finding.	§8 rule

5 New data requirements

Each requirement below must be satisfiable by the generator before its dependent anchors (§6) render. Priorities: **BLOCKER** the layout is dishonest without it; **HIGH** the layout degrades acceptably but loses its main advantage; **LATER** additive.

D-01 Per-source completeness BLOCKER T3

ANCHORS	<code>pick.completeness, source[n].completeness, show.setlist_confidence</code>
SHAPE	<code>{ songs_present: int, songs_total: int, missing: [song_position], partial: [song_position], basis: "tracklist" "runtime" null }</code>
SOURCE	Per-source file/track listings in the LB catalog, matched against the chronicle setlist. Requires exporting the track list, which the dossier payload currently discards.
METHOD	Normalise track titles (the catalog and the chronicle disagree on punctuation and subtitle handling — see <i>It's Alright, Ma (I'm Only Bleeding)</i> in the 1965 file), fuzzy-match to setlist positions, and record unmatched positions as missing. Where no track list exists, <code>basis</code> is <code>"runtime"</code> and only a coarse band may be shown.
ACCEPT	For 2010-03-29, all nine sources resolve with <code>basis: "tracklist"</code> . For 1965-06-01, LB-10364 and LB-12222 resolve to 1–2 songs of 12 and LB-08855 flags missing positions.
NULL	Render the runtime only; suppress the bar and the "18 / 18" figure entirely. Never infer completeness from duration alone in the verdict card.

D-02 Song-level performance history BLOCKER T3

ANCHORS	<code>song[n].premiere, song[n].gap, song[n].debut, stats.premiere_titles</code>
SHAPE	<code>{ tour_premiere: bool, career_debut: bool, last_played: date null, gap_shows: int null, times_played: int null }</code> per setlist position
SOURCE	A corpus-wide index of every setlist the install holds, keyed by song and date. The aggregate "2 new songs for this tour" already in the payload is a per-show summary with the identities stripped; the underlying comparison must be redone locally.
METHOD	For each song, find the previous performance within the tour (premiere if none) and within the career (debut if none); compute the show gap.
ACCEPT	The count of <code>tour_premiere: true</code> equals the payload's aggregate premiere count for every show in a 200-show sample. A mismatch is a hard failure, not a warning — it means the corpus index and the chronicle disagree.
NULL	Render the aggregate count in the stats rail with no per-song badges. Do <i>not</i> distribute the count across guessed songs.

D-03 Official release status BLOCKER T3

ANCHORS	<code>show.official_release, song[n].official_release</code>
SHAPE	<code>{ status: "full" "partial" "none", releases: [{title, year, songs:[position]}] }</code>
SOURCE	A release index (Bootleg Series and official live albums) mapped to show dates and song positions. No such source is ingested today; it must be acquired or hand-built.
METHOD	Date match for whole-show releases; date + title match for individual tracks.
ACCEPT	1975-12-08 (Madison Square Garden) must resolve non-empty — it is the highest-value case in the sample set and the reason the field is worth adding. A field that returns "none" for a show with a known official release is worse than no field.
NULL	Omit the row. The mock's "none indexed" is only permissible once the index exists and the lookup genuinely returned empty; absent the index, the row must not appear at all.

D-04 Cross-show rotation comparison HIGH T3

ANCHORS	<code>stats.rotation_rank, stats.rotation_context</code>
SHAPE	<code>{ pct: float, rank_in_run: int, run_size: int, run_median_pct: float }</code>
SOURCE	The per-show rotation percentage for every show in the same run/tour. Each show's own figure is already in its payload; the comparison requires reading its siblings.
ACCEPT	For 2010-03-29 the claim "biggest rotation of the run" either verifies against the other six Zepp nights or the phrase is replaced by the bare percentage.
NULL	State the percentage and the changed-song count with no superlative.

D-05 Generation classification HIGH T3

ANCHORS	<code>source[n].generation</code> , <code>pick.generation</code> , sources sort order
SHAPE	<code>{ class: "master" "low_gen" "silver" "vinyl" "broadcast" "unknown", evidence: string, confidence: "stated" "inferred" }</code>
SOURCE	Classifier over the lineage string plus the existing taper-propagation flag. The vocabulary is already rich enough to support this: mic model → preamp → recorder implies a master; "silver disc", "silver CD", a label name and catalogue number imply a commercial bootleg; "Near Mint Vinyl LPs" implies vinyl; "TV broadcast" / "PA recording" imply broadcast.
METHOD	Rule-based, not statistical, so <code>evidence</code> can quote the substring that fired. Anything unmatched is <code>unknown</code> — the class must never be guessed from rating or grade.
ACCEPT	Every classification carries a quotable <code>evidence</code> substring. Sample audit: LB-15005 → silver (label + catalogue number), LB-09493 (1965) → vinyl, LB-06654 (1965) → silver via the "BOOTLEG:" prefix, LB-08637 → unknown (no lineage on file).
NULL	Omit the tag. <code>confidence: "inferred"</code> renders with the §8 marker.

D-06 Venue enrichment LATER T3

ANCHORS	<code>venue.district</code> , <code>venue.capacity</code> , <code>venue.status</code>
SOURCE	External. Not available from setlist.fm coordinates alone.
NULL	City and country only. Drop "Odaiba" from the mock until this exists.

D-07 Structured run & tour position HIGH T3

ANCHORS	<code>run.*</code> (whole group), <code>context.chronicle</code> , <code>venue.run_nights</code>
SHAPE	<code>{ venue_run: {position, size, dates:[], dossier_urls:[]}, tour: {name, position, size, is_first, is_last}, city_history: [{year, venu</code>
SOURCE	Replaces the free-text "Other Bob Dylan shows in <city>" blob. The blob contains the data but as prose; the generator must emit it structured and stop emitting the prose version.
METHOD	Parse date/venue pairs from the chronicle, group by year and venue, resolve each date to a sibling dossier URL where one exists in this install.
ACCEPT	The city-history counts sum to the stated total — the defect that shipped in the first mock revision, where a "19 shows" heading sat above a list summing to 45. Enforce as an assertion in the generator, not a review step.
NULL	Suppress the run strip and the venue-history panel. Never render a partial count.

D-08 Pairwise source comparison HIGH T2

ANCHORS	<code>verdict.vs_runner_up</code> , <code>verdict.alternates[]</code>
SHAPE	<code>{ other_id, deltas: [{field, a, b, direction}], axis_wins: [{axis, id, value}] }</code>
SOURCE	Derivable from fields already present. Needs no new data — only the computation and the sentence templates.
METHOD	Diff the pick against rank 2 across rating, scan score, runtime, resolution, generation and character terms. Separately, find the leader on each axis (best raw scan, longest, highest resolution, only soundboard, only complete) and emit any leader that is not the pick as an alternate.
ACCEPT	For 2010-03-29: LB-08637 surfaces as best-scan alternate (90 vs 84) and LB-08493 as the 16/44.1 alternate. For 1965-06-01, where nine sources share identical rating, runtime and grade, the comparison must return <i>no meaningful deltas</i> and the anchor must collapse rather than manufacture a distinction.
NULL	Show the pick alone with no comparison prose.

D-09 Setlist completeness flag HIGH T3

ANCHORS	<code>show.setlist_confidence</code> , setlist section heading
SHAPE	<code>{ status: "complete" "partial" "unknown", evidence: string }</code>
SOURCE	Chronicle annotations plus a runtime sanity check: a 7-song setlist against 16 sources averaging full-concert length (1986-02-24) is almost certainly a truncated chronicle entry, not a 7-song concert.
ACCEPT	1986-02-24 flags <code>partial</code> . A show whose setlist song count implies a runtime wildly below its sources' median runtime must never be presented as a complete setlist.
NULL	Heading reads "Setlist · n songs listed" rather than "· n songs".

D-10 Bobtalk anchored to songs LATER T2

ANCHORS	<code>song[n].bobtalk</code> , <code>context.bobtalk_rest</code> , <code>verdict.pull_quote</code>
SOURCE	The existing bobtalk blob already carries positional cues — "(before Trust Yourself)", "(after Clean-Cut Kid)", "(plays Masters Of War)" in 1986-02-24.
METHOD	Parse the parenthetical cue, resolve the song title to a setlist position, attach as before/after. Unresolved lines stay in the context block.
ACCEPT	1986-02-24 attaches at least three lines to positions; 2010-03-29 attaches none and the anchor collapses without leaving an empty container.
NULL	Whole bobtalk blob renders in Context as today.

D-11 Format & file metadata LATER T3

ANCHORS	<code>source[n].resolution</code> , <code>source[n].filesize</code> , <code>source[n].disc_count</code> , <code>source[n].filecount</code>
SOURCE	Catalog file metadata. Resolution is <code>T2</code> where the lineage states it (2 of 9 sources for 2010-03-29); everything else requires the catalog's file records.
ACCEPT	No size figure renders unless it came from a file record. The first mock revision stated "~1.4 GB less" with no backing; that class of value is prohibited.
NULL	"—" in the Format column.

D-12 Catalog freshness LATER T3

ANCHORS	<code>source[n].added</code> , <code>show.newest_source_date</code>
SOURCE	Catalog date-added field, if retained on import.
NULL	Suppress the "new since" marker.

D-13 Medium classification (audio / video) HIGH T3

ANCHORS	<code>source[n].medium</code> , <code>show.media_available</code> , all runtime and completeness labels
SHAPE	<code>{ medium: "audio" "video" "audio_from_video", evidence: string }</code>
SOURCE	Lineage and note strings. The 1965-06-01 show is a <i>television broadcast</i> — its payload says "Mono TV broadcast" and "line recording" — yet the template's entire vocabulary ("transfer", "taper", "audience") assumes audio.
ACCEPT	1965-06-01 sources classify as broadcast-derived and the verdict card's language adapts.
NULL	Assume audio, as today.

6 Anchor reference (exhaustive)

Every anchor in the redesigned layout, in DOM order. Anchors are addressed in markup with `data-lb="{key}"` so coverage can be linted: the generator asserts that every emitted field has a matching anchor and that every anchor with `tier ≥ 3` has a defined fallback. *Fallback* describes what renders when the field is null.

6.1 Header

ANCHOR	TIER	FIELD / DERIVATION	FALLBACK
show.date.long	T1	Show date, spelled long. H1.	Required; no dossier without it
show.date.dow	T1	Weekday name.	Omit
show.venue	T1	Venue name. Subtitle line.	"Venue unknown"
show.city / .country	T1	City, country. Drops the region code the baseline shows ("Tokyo, 40 · Japan").	Omit clause
show.tour	T1	Tour name; may be a session label ("1965 Recording sessions & concerts").	Omit
show.net_number	T1	NET sequence number. Absent for pre-1988 shows.	Omit
show.year_index	T1	"2010 concert #14".	Omit
show.band_index	T2	"concert #47 with the 21st NET band" — parsed from the lineup blob's leading clause.	Omit
run.label	T3	D-07. "Zepp Tokyo run".	Suppress strip
run.position	T3	D-07. "night 7 of 7".	Suppress strip
run.dates[]	T3	D-07. One chip per night; current night non-interactive.	Suppress strip
run.dates[].url	T3	D-07. Sibling dossier URL, only where the install holds that show.	Render chip unlinked

6.2 Verdict

ANCHOR	TIER	FIELD / DERIVATION	FALLBACK
pick.lb_id	T1	Winning LB number.	Suppress §2 entirely
pick.url	T1	Detail-page URL; used by both the ID and the CTA.	Plain text ID
pick.taper.name	T1	Taper handle.	Omit taper line
pick.taper.confidence	T1	confirmed / propagated / absent. Drives the ✓ mark and its tooltip.	No mark
pick.medium	T3	D-13. Governs the noun used throughout ("transfer" vs "broadcast copy").	Assume audio
pick.type	T1	Audience / Soundboard chip.	Omit
pick.lineage_short	T1	Lineage string, truncated at the recorder.	Omit line
pick.lb_rating	T1	Catalog rating.	"_"
pick.scan_grade	T1	Local scan grade + /100 score. Absent for whole eras (all 18 sources in 1965-06-01).	Omit row
pick.resolution	T2	D-11. Parsed from lineage; catalog file records later.	Omit row
pick.generation	T3	D-05.	Omit row
pick.runtime	T2	Total plus per-set split.	Omit row
pick.completeness	T3	D-01. Bar + "n / total".	Omit row and bar
pick.curated_in[]	T1	Curator list names.	Omit row
show.official_release	T3	D-03.	Omit row (not "none")
verdict.why	T4	Composed from the ledger's positive terms plus the character string. Templated, never free-written.	Omit column
verdict.vs_runner_up	T2	D-08. Collapses when deltas are empty.	Omit column
verdict.alternates[]	T2	D-08. One clause per axis leader that is not the pick.	Omit column
ledger[]	T1	Score components, label + signed points.	Suppress disclosure
ledger.total	T1	Composite score, shown on the closed summary.	Suppress disclosure

6.3 Setlist

ANCHOR	TIER	FIELD / DERIVATION	FALLBACK
setlist.count	T1	Song count in the heading.	Omit heading suffix
show.setlist_confidence	T3	D-09. Changes the heading to "n songs listed" and adds a partial notice.	Assume complete
set[].label	T1	Set banding: main set / encore / broadcast segment. 1965 needs a per-broadcast band, which is where the repeated per-song note belongs (§7.3).	Single unbanded list
song[].position	T1	Ordinal.	Required
song[].title	T1	Song title. Subtitles rejoined with parentheses rather than emitted as a second line.	Required
song[].writers	T2	Parsed from the parenthetical credit; shown only for non-Dylan or co-written songs.	Omit
song[].instruments[]	T2	Expanded from the lineup blob's song-range notation ("2, 4-7, 10, 14, 18 harmonica"). Dylan's own instrument and named sidemen only.	Omit chips
song[].premiere	T3	D-02. "new to tour" / "debut" badge.	No badge
song[].gap	T3	D-02. "first since 1994" style marker.	No badge
song[].bobtalk	T2	D-10. Inline quote above or below the row.	Stays in Context
song[].official_release	T3	D-03. Marker where this performance is officially released.	No marker
stats.rotation	T1	Changed-song count and percentage vs previous show.	Omit stat
stats.rotation_rank	T3	D-04. The superlative clause only.	Percentage alone
stats.premiere_count	T1	Aggregate count.	Omit stat
stats.premiere_titles	T3	D-02. Named titles.	Count without names
stats.instrument_tally	T2	"harp on 8 of 18" and the guitar/keyboard split.	Omit stat
band.members[]	T2	Parsed from the lineup blob into name + instruments. Dylan first, then as listed.	Raw lineup line in Context
band.label	T2	"21st NET band" or "solo".	"Personnel"

6.4 Sources

ANCHOR	TIER	FIELD / DERIVATION	FALLBACK
<code>sources.count</code>	T1	Total circulating sources.	Required
<code>sources.tape_count</code>	T1	Independent tapes, in the heading credit — more meaningful than the family count, which counts solos as families.	Omit clause
<code>sources.visible_n</code>	T2	Rows shown before the collapse. Rule in §8.3.	Show all
<code>source[.].lb_id / .url</code>	T1	Identifier and link.	Required
<code>source[.].taper</code>	T1	Handle + confidence mark.	Omit sub-line (never "unknown")
<code>source[.].generation</code>	T3	D-05. Tag beside the taper.	Omit tag
<code>source[.].medium</code>	T3	D-13.	Assume audio
<code>source[.].type</code>	T1	Audience / Soundboard. Column suppressed when empty for the whole show.	Suppress column
<code>source[.].lb_rating</code>	T1	Rating.	"_"
<code>source[.].scan_grade</code>	T1	Grade + score. Column suppressed when empty for the whole show (all of 1965-06-01).	Suppress column
<code>source[.].runtime</code>	T2	Total minutes; split in the tooltip.	"_"
<code>source[.].completeness</code>	T3	D-01. Also gates the fragment rule in §8.2.	Omit
<code>source[.].resolution</code>	T2	D-11.	"_"
<code>source[.].character</code>	T1	Character sentence, with the "Grade X (n/100). LBnnnn:" prefix and the redundant intra-family rank stripped.	Empty cell
<code>source[.].flags[]</code>	T1	Defects — "audible mains hum", "no lineage on file".	Omit
<code>source[.].lineage</code>	T1	Full string; moved out of rows into the section footnote.	Omit from footnote
<code>source[.].curated_in[]</code>	T1	Curator lists; shown as a mark, not a sentence per row.	Omit
<code>source[.].rank</code>	T1	Pick rank. Non-contiguous within the visible set once rows collapse — intentional, and the reason the column is labelled "Rank" not "#".	Omit column
<code>family[.].id / .label</code>	T1	Family name. Rendered only when size > 1 (§8.4). Label prefers the taper handle ("oscar's tape") over "Family A".	No band
<code>family[.].size</code>	T1	Member count.	No band
<code>family[.].confidence</code>	T1	Match percentage + bar; drives the low-confidence tag (§8.5).	Bar omitted
<code>family[.].basis</code>	T3	What the match was computed on — length, waveform, or note text. Needed to make the confidence figure interpretable; a bare "4%" is currently unreadable.	Omit tooltip

6.5 Context, Related, Provenance

ANCHOR	TIER	FIELD / DERIVATION	FALLBACK
context.chronicle	T4	One or two sentences composed from run.* , show.tour and the chronicle's own line. Replaces the raw blob. Must not assert tour-position facts absent from D-07.	Chronicle line verbatim
context.bobtalk	T1	Unattributed remainder after D-10.	"No bobtalk recorded" notice
context.session_notes	T1	Genuine session notes — broadcast dates, recording format, circulation remarks. Explicitly <i>not</i> the venue-history blob (§7.1).	Omit
venue.name / .city	T1	Location card head.	Omit card
venue.district	T3	D-06.	Omit
venue.coords	T1	Coordinates + attribution, as text. No map.	Omit line
venue.run_nights	T3	D-07.	Omit clause
venue.city_total	T3	D-07. Must equal the sum of city_history[] — asserted, not reviewed.	Omit clause
city_history[]	T3	D-07. Year + venue + count rows.	Suppress panel
xrefs[]	T1	Cross-reference links; deep URL where known, site root otherwise.	Omit missing rows
prov.credits	T1	Source attributions.	Required
prov.local_fields[]	T3	New: the list of fields on <i>this</i> page that came from local analysis, generated from the tier map so the disclosure is accurate per show rather than boilerplate.	Current blanket sentence
prov.stamps	T1	Generation time, master version, channel.	Required

7 Defects in the current generator

These are independent of the redesign and should be fixed regardless; several are the reason the redesign looked like it needed invented data.

7.1 Venue-history blob rendered as session notes

2010-03-29's `Notes` field contains forty lines of "Other Bob Dylan shows in Tokyo, Japan" and becomes the tallest element on the page. The data is valuable but belongs in `city_history[]` (D-07). Route it there and stop emitting it as prose.

7.2 Song notes mis-joined

In the same file, setlist rows 1–12 carry fragments of the venue-history blob ("March 1978; April 2014") as per-song notes, and row 13 — *Thunder On The Mountain* — carries the show-level rotation statistic ("new songs (72%) compared to previous concert"). Two different show-level strings are being distributed across song rows. Fix the join before D-02 lands, or per-song badges will inherit the same corruption.

7.3 Set-level notes repeated per song

1965-06-01 repeats "broadcast by BBC TV-1, 19 June 1965" on six consecutive rows and the 26 June equivalent on the next six. That is a set-level fact and should band the two broadcast segments via `set[].label`.

7.4 Fragments ranked against complete shows

1965-06-01 ranks LB-10364 (3 min) at #10 and LB-12222 (3 min) at #11, above complete 66-minute transfers, and LB-08855 (50 min of a 66-minute broadcast) at #9. A single-song excerpt is not a candidate for "the copy to get". See §8.2.

7.5 Columns and rows rendered empty

1965-06-01 renders Type and AI-grade columns that are empty for all eighteen rows, and two rows (LB-06654, LB-14937) with no rating, no runtime and no lineage but a 40.0 composite score. Suppress empty columns per show; give scoreless rows either an explanation or exclusion.

7.6 Always-on warning badges

The orange `review` badge appears on every tape family in all five sampled dossiers, including thirteen consecutive solo "families" in 1965-06-01. A warning that always fires conveys nothing. Reserve it for genuine low-confidence groupings (§8.5).

7.7 Solo families rendered as families

Ten of 1965-06-01's thirteen "families" have one member, each costing a full banner row reading "Solo · 1 source · one master tape · review". See §8.4.

7.8 Template whitespace

Roughly 40% of each output file is blank lines emitted by the template's control structures — several hundred lines before the masthead. Cosmetic, but it triples the page's byte size for no content.

8 Rendering rules & degradation

8.1 Null discipline

Absence of data is rendered as absence, never as a finding. "Taper unknown" is prohibited where the dossier simply has no taper field; the sub-line is omitted. "None indexed" is permitted only when an index was actually consulted and returned empty. Any anchor whose field is null renders the fallback in §6 and nothing else.

8.2 Fragments

A source whose completeness (D-01) is below a threshold — provisionally 60% of setlist positions, or a runtime below 60% of the median source runtime where completeness is unavailable — is excluded from the pick ranking, listed in a separate "Excerpts & fragments" group after the main table, and labelled with what it contains ("1 song · If You Gotta Go, Go Now"). It remains fully linked; it is simply not a candidate for the verdict.

8.3 Collapse threshold

Show the pick plus the next three ranked sources, then collapse. If a source outside the top four is an axis leader under D-08, promote it into the visible set. If the show has five or fewer sources, do not collapse at all (1990-01-25 has five).

8.4 Family banding

Render a family band only when the family has more than one member; single-source families render as plain rows with no banner. Prefer a human label derived from the taper handle. Sort within a family by generation (master → low gen → silver → vinyl) and then by rank, so the band reads as a lineage rather than a score list.

8.5 Confidence marking

A family confidence below 50% earns the low-confidence tag; at or above 50% the bar alone suffices. The tag's tooltip states the basis (D-05/family basis) so the reader knows what the percentage measured. Any inferred field renders with a dotted underline and a tooltip naming the inference — this is the mechanism that lets generation tags and composed sentences appear without overclaiming.

8.6 Sparse-show behaviour

1965-06-01 is the stress case: eighteen sources, nine sharing an identical rating, runtime and grade, no scan grades at all, no taper for most, and two three-minute fragments. Under these rules it renders as: no scan-grade column, no Type column, fragments split out, ten solo families as plain rows, one real family banded, D-08 returning no meaningful deltas so the verdict card shows the pick alone with its ledger. That is the correct outcome — the page should say "these are all much the same" rather than manufacture a hierarchy.

9 Open questions

1. Can the catalog export per-source track listings? D-01 and the fragment rule both depend on it, and it is the single highest-value addition in this spec.
2. Does the install hold enough setlist corpus to compute D-02 locally, or does song history need an external index?
3. Is there an acceptable source for official-release mapping (D-03), and is hand-building one for the Bootleg Series volumes plus official live albums proportionate?

4. Should the sibling-night strip link only to dossiers this install has generated, or to a stable URL pattern regardless of whether the page exists?
5. Confirm the fragment threshold in §8.2 against a wider sample before fixing it at 60%.
6. Does the "AI grade" label survive? "Scanned quality" is used throughout the mock and the spec, on the grounds that it describes what was measured rather than what measured it.

Rev A · 9 September 2026 · Sample set: 1965-06-01, 1975-12-08, 1986-02-24, 1990-01-25, 2010-03-29. Reference mock: *Show Dossier — Redesign (2010-03-29).html*. Requirement IDs D-01...D-13 are stable and may be cited in issue tracking.